




Gaussian Mixture Models, Expectation Maximization

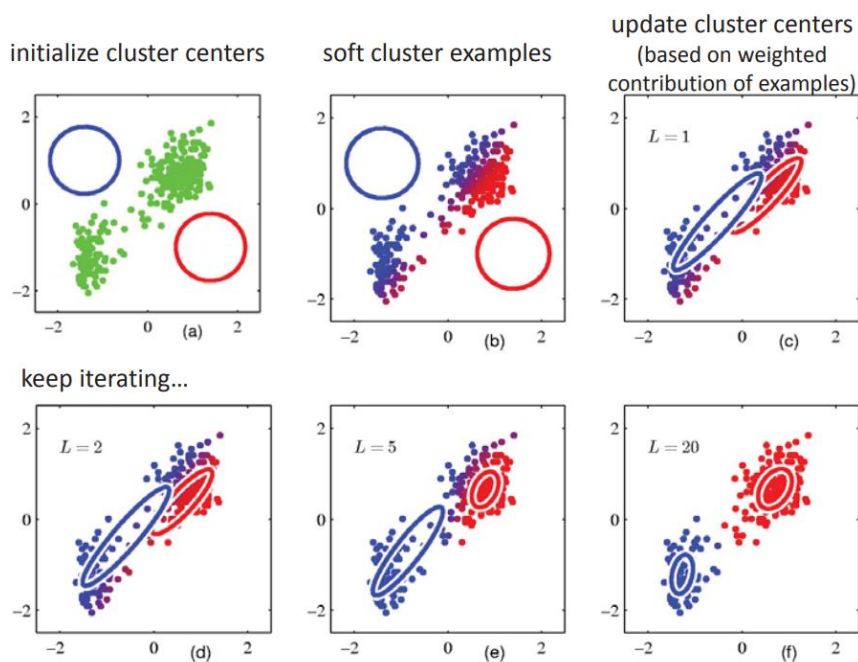
Gaussian Mixture Models

- Assume data came from **mixture of Gaussians** (elliptical data)
- Assign data to cluster with certain **probability** (soft clustering)



- Very similar at high-level to *k*-means: iterate between assigning examples and updating cluster centers

GMM Example




Univariate Gaussian Distribution

(scalar) random variable X

parameters: (scalar) mean μ , (scalar) variance σ^2

$$X \sim N(\mu, \sigma^2) \quad p(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$



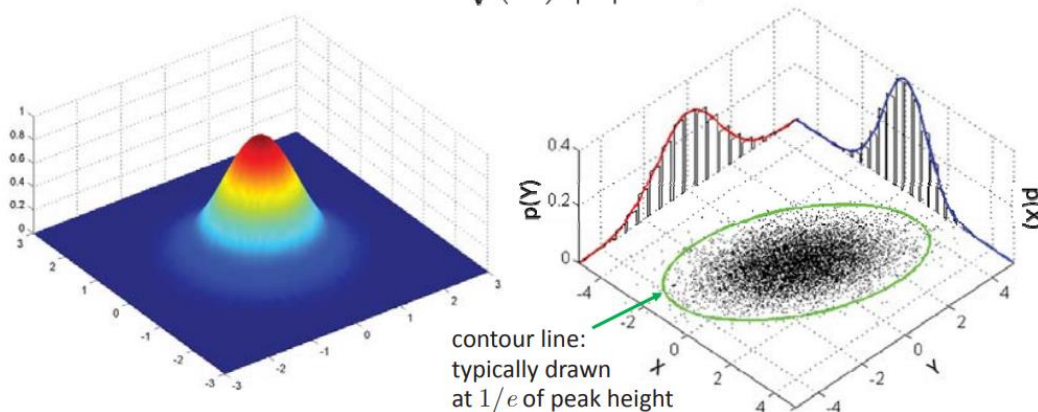
Multivariate Gaussian Distribution

random variable vector $\mathbf{X} = [X_1, \dots, X_n]^T$

parameters: mean vector $\boldsymbol{\mu} \in \mathbb{R}^n$

covariance matrix $\boldsymbol{\Sigma}$ (symmetric, positive definite)

$$\mathbf{X} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \quad p(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^n |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$



Covariance Matrix

Recall for pair of r.v.'s X and Y , covariance is defined as

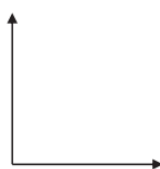
$$\text{cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

For $\mathbf{X} = [X_1, \dots, X_n]^T$, **covariance matrix** summarizes covariances across all pairs of variables:

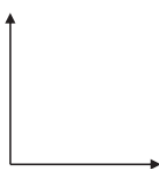
$$\Sigma = \mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^T]$$

Σ is $n \times n$ matrix s.t. $\Sigma_{ij} = \text{cov}(X_i, X_j)$

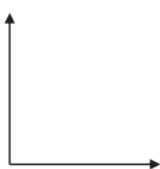
parameters

$$\Sigma = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \dots & \sigma_{1n} \\ \sigma_{21} & \sigma_2^2 & \dots & \sigma_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{n1} & \sigma_{n2} & \dots & \sigma_n^2 \end{bmatrix}$$


params

$$\Sigma = \begin{bmatrix} \sigma_1^2 & & 0 \\ & \sigma_2^2 & \\ 0 & & \ddots \\ & & & \sigma_n^2 \end{bmatrix}$$



params

$$\Sigma = \begin{bmatrix} \sigma^2 & & 0 \\ & \sigma^2 & \\ 0 & & \ddots \\ & & & \sigma^2 \end{bmatrix}$$




GMMs as Generative Model

- There are k components
- Component j
 - has associated mean vector μ_j and covariance matrix Σ_j
 - generates data from $N(\mu_j, \Sigma_j)$
- Each example $\mathbf{x}^{(i)}$ is generated according to following recipe:
 - pick component j at random with probability ϕ_j
 - sample $\mathbf{x}^{(i)} \sim N(\mu_j, \Sigma_j)$



GMM Optimization

Assume supervised setting (known cluster assignments)

MLE for univariate Gaussian

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x^{(i)} \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n \left(x^{(i)} - \hat{\mu} \right)^2$$

sum over points generated
from this Gaussian

MLE for multivariate Gaussian

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \quad \hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n \left(\mathbf{x}^{(i)} - \hat{\mu} \right) \left(\mathbf{x}^{(i)} - \hat{\mu} \right)^T$$

Expectation Maximization

- Clever method for maximizing marginal likelihoods

$$\arg \max_{\theta} \prod_{i=1}^n P \left(\mathbf{x}^{(i)} \right) = \arg \max_{\theta} \prod_{i=1}^n \sum_{j=1}^k P \left(\mathbf{x}^{(i)}, z^{(i)} = j \right)$$

- Excellent approach for unsupervised learning
- Can do “trivial” things (upcoming example)
- One of most general unsupervised approaches with many other uses (e.g. HMM inference)

Overview

- Begin with guess for model parameters
- Repeat until convergence
 - Update latent variables based on our expectations [E-step]
 - Update model parameters to maximize log likelihood [M-step]

Silly Example

Let events be “grades in a class”

component 1 = gets an A $P(A) = \frac{1}{2}$

component 2 = gets a B $P(B) = p$

component 3 = gets a C $P(C) = 2p$

component 4 = gets a D $P(D) = \frac{1}{2} - 3p$ (note $0 \leq p \leq 1/6$)

Assume we want to estimate p from data. In a given class, there were

a A's, b B's, c C's, d D's.

What is the MLE of p given a, b, c, d ?

so if class got

a	b	c	d
14	6	9	10

Same Problem with Hidden Information

Someone tells us that

of high grades (A's + B's) = h

of C's = c

of D's = d

What is the MLE of p now?

Remember

$P(A) = \frac{1}{2}$

$P(B) = p$

$P(C) = 2p$

$P(D) = \frac{1}{2} - 3p$

We can answer this question circularly:

EXPECTATION

If we know value of p ,
we could compute **expected** values of a and b .

MAXIMIZATION

If we know expected values of a and b ,
we could compute **maximum** likelihood value of p .

EM for Our Silly Example

- Begin with initial guess for p
- Iterate between Expectation and Maximization to improve our estimates of p and a & b

- Define $p^{(t)}$ = estimate of p on t^{th} iteration
 $b^{(t)}$ = estimate of b on t^{th} iteration

- Repeat until convergence

$$\boxed{\text{E-step}} \quad b^{(t)} = \frac{p^{(t)}}{\frac{1}{2} + p^{(t)}} h = \mathbb{E}[b|p^{(t)}]$$

$$\boxed{\text{M-step}} \quad p^{(t+1)} = \frac{b^{(t)} + c}{6(b^{(t)} + c + d)} = \text{MLE of } p \text{ given } b^{(t)}$$

EM Convergence

- **Good news:** converging to local optima is guaranteed
- **Bad news:** local optima

Aside (idea behind **convergence proof**)

- likelihood must increase or remain same between each iteration [not obvious]
- likelihood can never exceed 1 [obvious]
- so likelihood must converge [obvious]

In our example, suppose we had

$$h = 20, c = 10, d = 10$$

$$p^{(0)} = 0$$

Error generally decreases by constant factor each time step
(e.g. convergence is linear)

t	$p^{(t)}$	$b^{(t)}$
0	0	0
1	0.0833	2.857
2	0.0937	3.158
3	0.0947	3.185
4	0.0948	3.187
5	0.0948	3.187
6	0.0948	3.187

Final Comments

EM is not magic

- Still optimizing non-convex function with lots of local optima
- Computations are just easier (often, significantly so!)

Problems

- EM susceptible to local optima

⇒ reinitialize at several different initial parameters

Extensions

- EM looks at maximum log likelihood of data

⇒ also possible to look at maximum *a posteriori*

Clustering methods: Comparison

	Hierarchical	K-means	GMM
Running time	naively, $O(N^3)$	fastest (each iteration is linear)	fast (each iteration is linear)
Assumptions	requires a similarity / distance measure	strong assumptions	strongest assumptions
Input parameters	none	K (number of clusters)	K (number of clusters)
Clusters	subjective (only a tree is returned)	exactly K clusters	exactly K clusters