

q-learning-bitcoins

December 21, 2023

```
[65]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
```

```
[66]: btc_data = pd.read_csv('/kaggle/input/binance-top-cryptocurrencies/BTC.csv')
btc_data.head()
```

```
[66]:
```

	unix	date	symbol	open	high	low	\
0	1.610496e+12	2021-01-13 00:00:00	BTC/USDT	34049.15	34049.15	33589.57	
1	1.610410e+12	2021-01-12 00:00:00	BTC/USDT	35410.37	36628.00	32531.00	
2	1.610323e+12	2021-01-11 00:00:00	BTC/USDT	38150.02	38264.74	30420.00	
3	1.610237e+12	2021-01-10 00:00:00	BTC/USDT	40088.22	41350.00	35111.11	
4	1.610150e+12	2021-01-09 00:00:00	BTC/USDT	40586.96	41380.00	38720.00	

	close	Volume BTC	Volume USDT	trade count
0	33606.99	655.275245	2.215246e+07	11838.0
1	34051.24	133948.151996	4.651302e+09	2674145.0
2	35404.47	249131.539943	8.426880e+09	4431451.0
3	38150.02	118209.544503	4.604035e+09	2628050.0
4	40088.22	75785.979675	3.054779e+09	1998156.0

```
[67]: btc_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1247 entries, 0 to 1246
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   unix            1247 non-null  float64
1   date            1247 non-null  object
2   symbol          1247 non-null  object
3   open            1247 non-null  float64
4   high            1247 non-null  float64
5   low             1247 non-null  float64
6   close           1247 non-null  float64
7   Volume BTC      1247 non-null  float64
```

```

      8   Volume USDT   1247 non-null   float64
      9   tradecount   1123 non-null   float64
dtypes: float64(8), object(2)
memory usage: 97.5+ KB

```

```
[68]: btc_data.describe()
```

```
[68]:
```

	unix	open	high	low	close \
count	1.247000e+03	1247.000000	1247.000000	1247.000000	1247.000000
mean	1.406850e+12	8819.920313	9085.641099	8518.824587	8843.271716
std	4.679284e+11	4586.887774	4794.111621	4340.497940	4638.470396
min	1.502928e+09	3189.020000	3276.500000	2817.000000	3189.020000
25%	1.529755e+12	6387.980000	6532.575000	6253.105000	6389.485000
50%	1.556669e+12	8175.630000	8348.620000	7917.000000	8175.640000
75%	1.583582e+12	10195.095000	10438.500000	9869.230000	10208.395000
max	1.610496e+12	40586.960000	41950.000000	38720.000000	40582.810000

	Volume BTC	Volume USDT	tradecount
count	1247.000000	1.247000e+03	1.123000e+03
mean	45316.815018	4.528942e+08	5.037790e+05
std	34754.098993	6.010777e+08	4.463849e+05
min	225.760000	9.677533e+05	1.183800e+04
25%	25262.445595	1.697613e+08	2.156135e+05
50%	38815.409893	3.048624e+08	3.561340e+05
75%	58523.946985	5.148775e+08	6.378765e+05
max	402201.673764	8.426880e+09	4.431451e+06

```
[69]: btc_data.isnull().sum()
```

```
[69]:
```

unix	0
date	0
symbol	0
open	0
high	0
low	0
close	0
Volume BTC	0
Volume USDT	0
tradecount	124

dtype: int64

```
[70]: # Handle missing values in 'tradecount' column
btc_data['tradecount'].fillna(btc_data['tradecount'].mean(), inplace=True)
btc_data.isnull().sum()
```

```
[70]:
```

unix	0
date	0

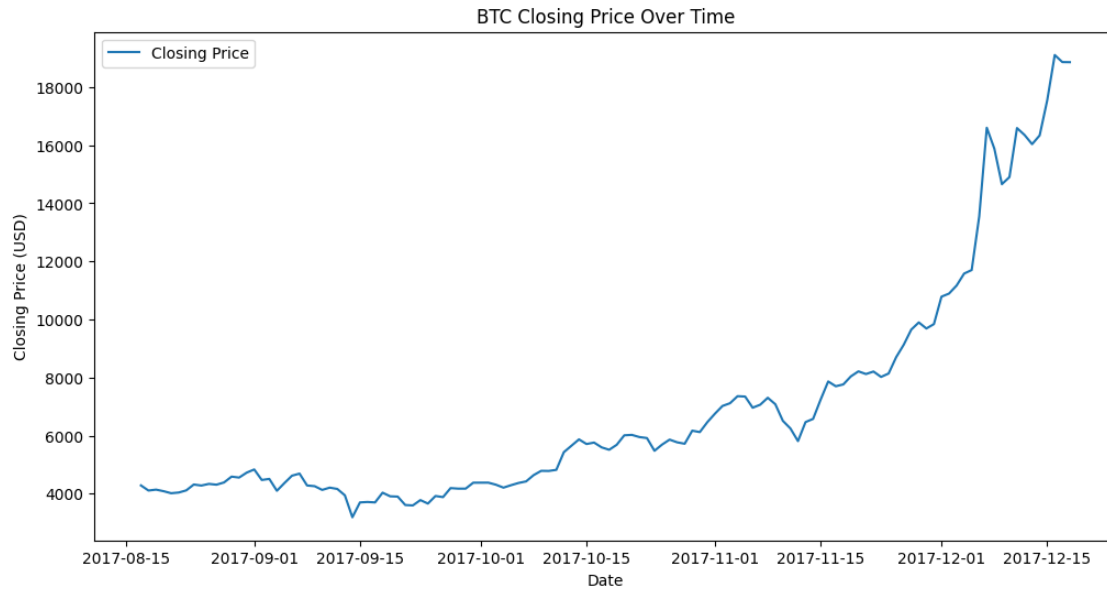
```
symbol      0
open        0
high        0
low         0
close       0
Volume BTC  0
Volume USDT 0
tradecount  0
dtype: int64
```

```
[71]: btc_data.dtypes
```

```
[71]: unix          float64
date            object
symbol         object
open           float64
high           float64
low            float64
close          float64
Volume BTC     float64
Volume USDT    float64
tradecount     float64
dtype: object
```


```
[72]: # convert date format
btc_data['date'] = pd.to_datetime(btc_data['date'], format="%Y-%m-%d",
    ↪errors='coerce')
```

```
[73]: # Visualize data
# Time series plot
plt.figure(figsize=(12, 6))
plt.plot(btc_data['date'], btc_data['close'], label='Closing Price')
plt.title('BTC Closing Price Over Time')
plt.xlabel('Date')
plt.ylabel('Closing Price (USD)')
plt.legend()
plt.show()
```



[74]: *#Plot the trading volume over time.*

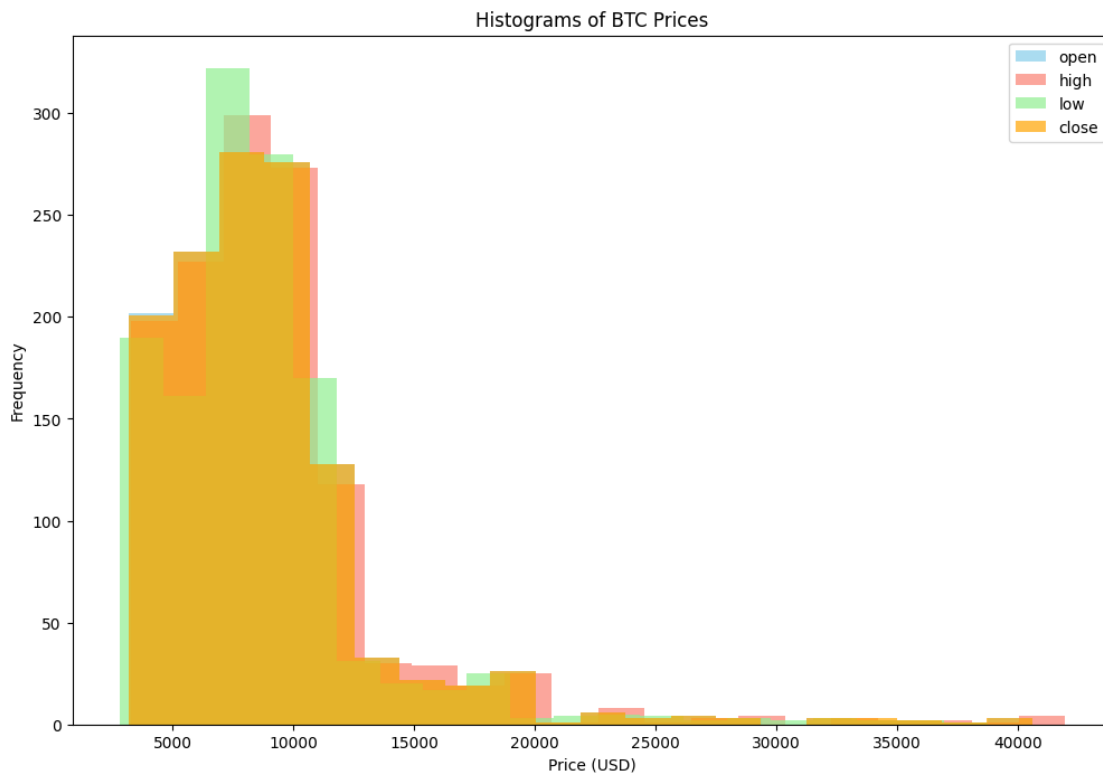
```
plt.figure(figsize=(12, 6))
plt.bar(btc_data['date'], btc_data['Volume BTC'], color='green', alpha=0.5)
plt.title('BTC Trading Volume Over Time')
plt.xlabel('Date')
plt.ylabel('Volume BTC')
plt.show()
```



```
[75]: #Histograms
colors = ['skyblue', 'salmon', 'lightgreen', 'orange']
plt.figure(figsize=(12, 8))

for i, column in enumerate(['open', 'high', 'low', 'close']):
    plt.hist(btc_data[column], bins=20, alpha=0.7, color=colors[i],
            label=column)


plt.title('Histograms of BTC Prices')
plt.xlabel('Price (USD)')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



```
[76]: # Select only numeric columns for correlation matrix
num_col = btc_data.select_dtypes(include=['float64']).columns
corr = btc_data[num_col].corr()
```

```
[77]: # Plot the correlation heatmap
import seaborn as sns
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```



```
[78]: # Feature scaling using Min-Max scaling
scaler = MinMaxScaler()
fs = ['open', 'high', 'low', 'close', 'Volume BTC', 'Volume USDT', 'tradecount']

# Apply Min-Max scaling to selected features
btc_data[fs] = scaler.fit_transform(btc_data[fs])
```

```
[79]: # Feature engineering
window_size = 10
btc_data['SMA'] = btc_data['close'].rolling(window=window_size).mean()
```

```
[80]: # Display the preprocessed data
print("Preprocessed Data:")
print(btc_data.head())
```

Preprocessed Data:

	unix	date	symbol	open	high	low	close	\
0	1.610496e+12	NaT	BTC/USDT	0.825183	0.795704	0.857103	0.813450	
1	1.610410e+12	NaT	BTC/USDT	0.861581	0.862386	0.827619	0.825330	
2	1.610323e+12	NaT	BTC/USDT	0.934838	0.904708	0.768822	0.861519	
3	1.610237e+12	NaT	BTC/USDT	0.986664	0.984486	0.899482	0.934941	
4	1.610150e+12	NaT	BTC/USDT	1.000000	0.985261	1.000000	0.986773	

	Volume BTC	Volume USDT	trade count	SMA
0	0.001069	0.002514	0.000000	NaN
1	0.332663	0.551909	0.602385	NaN
2	0.619206	1.000000	1.000000	NaN
3	0.293510	0.546299	0.591955	NaN
4	0.187972	0.362431	0.449433	NaN

1 Define Q-learning parameters

```
[81]: gamma = 0.9 # Discount factor
alpha = 0.1 # Learning rate
epsilon = 0.1 # Exploration-exploitation trade-off
```

2 Define Q-learning Algorithm

```
[82]: def q_learning(Q, state, action, reward, next_state):
    current_value = Q[state, action]
    max_future_value = np.max(Q[next_state, :])
    new_value = (1 - alpha) * current_value + alpha * (reward + gamma *
    ↪ max_future_value)
    Q[state, action] = new_value
    return Q
```

```
[83]: train_data, test_data = train_test_split(btc_data, test_size=0.2,
    ↪ random_state=42)
```

```
[84]: # Initialize Q-table
num_states = 100
num_actions = 2

Q = np.zeros((num_states, num_actions))
```

3 Training the Q-learning model

```
[85]: for index, row in train_data.iterrows():
        state = int(index % num_states)
        action = 1 if row['close'] < row['SMA'] else 0 # Buy if the price is below
        ↪ the moving average, else sell
        reward = row['close'] - row['SMA'] # Reward is the difference from the
        ↪ moving average
        next_state = int((index + 1) % num_states) # Simple next state
        ↪ representation, adjust accordingly

        Q = q_learning(Q, state, action, reward, next_state)
```

4 Testing the Q-learning model

```
[86]: # Testing the Q-learning model
total_reward = 0
actions_taken = []

for index, row in test_data.iterrows():
    state = int(index % num_states)
    action = np.argmax(Q[state, :])
    actions_taken.append(action)
    reward = row['close'] - row['SMA']
    total_reward += reward
```

5 Visualize actions taken during testing

```
[87]: plt.figure(figsize=(8, 6))
plt.plot(test_data['date'], test_data['close'], label='Closing Price')


# Filter buy signals
buy_signals = test_data['date'][np.array(actions_taken) == 1]
plt.scatter(buy_signals, test_data.loc[test_data['date'].
    ↪ isin(buy_signals)]['close'],
            color='green', label='Buy Signal', marker='^')

# Filter sell signals
sell_signals = test_data['date'][np.array(actions_taken) == 0]
plt.scatter(sell_signals, test_data.loc[test_data['date'].
    ↪ isin(sell_signals)]['close'],
            color='red', label='Sell Signal', marker='v')

plt.title('BTC Closing Price and Trading Signals')
plt.xlabel('Date')
```



```
plt.ylabel('Closing Price (USD)')
plt.legend()
plt.show()
```



```
[88]: # Visualize actions taken during testing with annotations
plt.figure(figsize=(8, 6))
plt.plot(test_data['date'], test_data['close'], label='Closing Price')

# Filter buy signals
buy_signals = test_data['date'][np.array(actions_taken) == 1]
plt.scatter(buy_signals, test_data['close'].loc[test_data['date']
    ↪isin(buy_signals)],
            color='green', label='Buy Signal', marker='^')

# Filter sell signals
sell_signals = test_data['date'][np.array(actions_taken) == 0]
plt.scatter(sell_signals, test_data['close'].loc[test_data['date']
    ↪isin(sell_signals)],
            color='red', label='Sell Signal', marker='v')
```

```

# Annotate Buy signals
for date, price in zip(buy_signals, test_data['close'].loc[test_data['date']
    ↳isin(buy_signals)]):
    plt.annotate('Buy', (date, price), textcoords="offset points", xytext=(0,
    ↳10), ha='center', fontsize=8, color='green')

# Annotate Sell signals
for date, price in zip(sell_signals, test_data['close'].loc[test_data['date']
    ↳isin(sell_signals)]):
    plt.annotate('Sell', (date, price), textcoords="offset points", xytext=(0,
    ↳-10), ha='center', fontsize=8, color='red')

plt.title('BTC Closing Price and Trading Signals with Annotations')
plt.xlabel('Date')
plt.ylabel('Closing Price (USD)')
plt.legend()
plt.show()

```

