

# auto-regressive-methods

December 29, 2023

```
[178]: import pandas as pd
import numpy as np
import seaborn as sns

import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
from sklearn.metrics import mean_squared_error
```

```
[179]: data = pd.read_csv('/kaggle/input/air-passengers/AirPassengers.csv')
data.head()
```

```
[179]:      Month  #Passengers
0  1949-01           112
1  1949-02           118
2  1949-03           132
3  1949-04           129
4  1949-05           121
```

```
[180]: data = data.rename(columns={"#Passengers": "Passengers"}, inplace=False)
data.head()
```

```
[180]:      Month  Passengers
0  1949-01           112
1  1949-02           118
2  1949-03           132
3  1949-04           129
4  1949-05           121
```


```
[181]: data.columns = ['Month', 'Passengers']
data['Month'] = pd.to_datetime(data['Month'], format='%Y-%m')
data = data.set_index('Month')
data.head()
```

```
[181]:      Passengers
Month
1949-01-01      112
1949-02-01      118
```

1949-03-01	132
1949-04-01	129
1949-05-01	121

## 1 Plot time series data

```
[182]: data.plot(y='Passengers', figsize=(20, 4), color='blue', linestyle='-',
    ↪marker='o', markersize=5, label='Passenger Traffic')
plt.grid(True)
plt.legend(loc='best')
plt.title('Airline Passenger Traffic Over Time')
plt.xlabel('Date')
plt.ylabel('Number of Passengers')
plt.show(block=False)
```

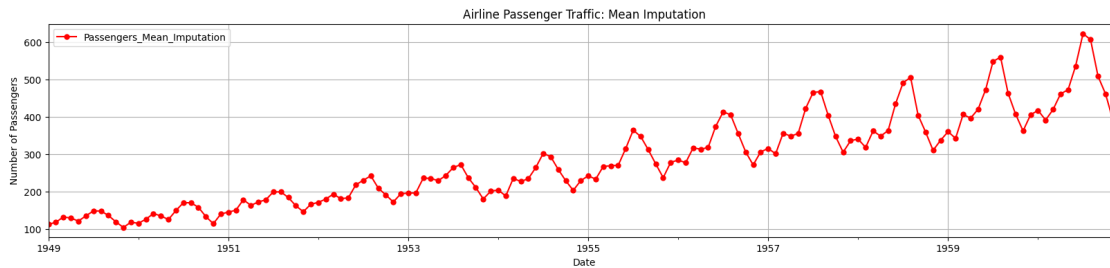


## 2 Missing value treatment

### 2.1 Mean imputation

```
[183]: data['Passengers_Mean_Imputation'] = data['Passengers'].
    ↪fillna(data['Passengers'].mean())
```

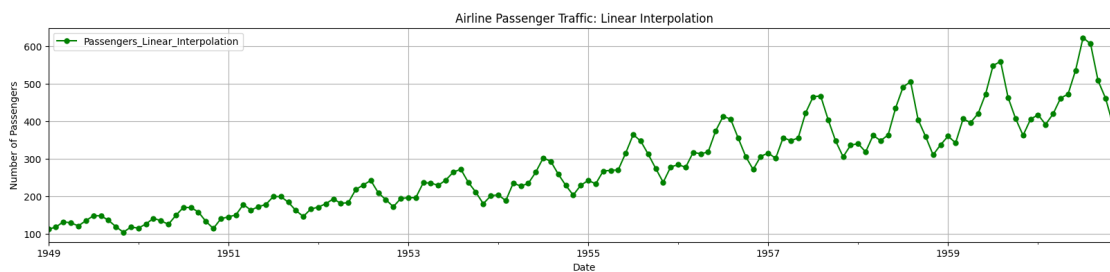
```
[184]: data[['Passengers_Mean_Imputation']].plot(figsize=(20, 4), grid=True,
    ↪legend=True, color='red', linestyle='-', marker='o', markersize=5)
plt.title('Airline Passenger Traffic: Mean Imputation')
plt.xlabel('Date')
plt.ylabel('Number of Passengers')
plt.show(block=False)
```



## 2.2 Linear interpolation

```
[185]: data['Passengers_Linear_Interpolation'] = data['Passengers'].
        ↪ interpolate(method='linear')
```

```
[186]: data[['Passengers_Linear_Interpolation']].plot(figsize=(20, 4), grid=True,
        ↪ legend=True, color='green', linestyle='-', marker='o', markersize=5)
plt.title('Airline Passenger Traffic: Linear Interpolation')
plt.xlabel('Date')
plt.ylabel('Number of Passengers')
plt.show(block=False)
```



```
[187]: data.head()
```

```
[187]:
```

	Passengers	Passengers_Mean_Imputation	\
Month			
1949-01-01	112	112	
1949-02-01	118	118	
1949-03-01	132	132	
1949-04-01	129	129	
1949-05-01	121	121	

	Passengers_Linear_Interpolation
Month	
1949-01-01	112

1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

```
[188]: data['Passengers'] = data['Passengers_Linear_Interpolation']
data.
↳drop(columns=['Passengers_Mean_Imputation', 'Passengers_Linear_Interpolation'], inplace=True)
data.head()
```

```
[188]:
```

Month	Passengers
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

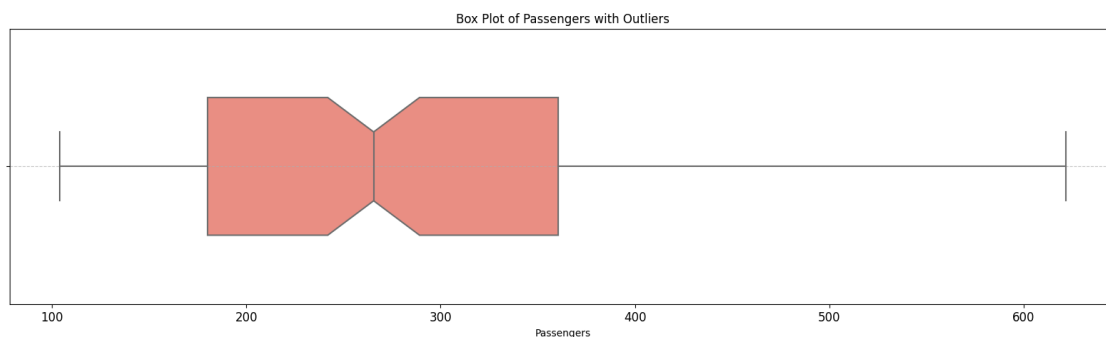
### 3 Outlier detection

#### 3.1 Box plot and interquartile range

```
[189]: import seaborn as sns
plt.figure(figsize=(20, 5))

sns.boxplot(x=data['Passengers'], color='salmon', width=0.5, notch=True)

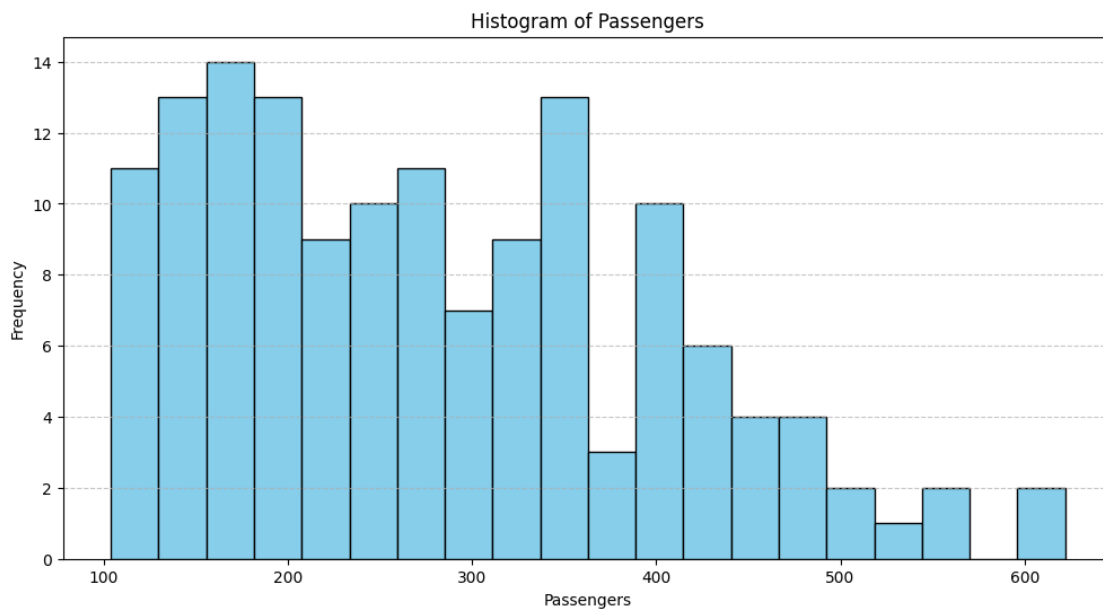
plt.title('Box Plot of Passengers with Outliers')
plt.xlabel('Passengers')
plt.xticks(fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



## 3.2 Histogram plot

```
[190]: plt.figure(figsize=(12, 6))
plt.hist(data['Passengers'], bins=20, color='skyblue', edgecolor='black')

plt.title('Histogram of Passengers')
plt.xlabel('Passengers')
plt.ylabel('Frequency')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



## 4 Time series Decomposition


### 4.1 Additive seasonal decomposition

```
[191]: from statsmodels.tsa.seasonal import seasonal_decompose

# Perform decomposition
result = seasonal_decompose(data['Passengers'], model='additive')
fig, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(10, 8), sharex=True)

result.observed.plot(ax=ax1, title='Observed')
result.trend.plot(ax=ax2, title='Trend')
result.seasonal.plot(ax=ax3, title='Seasonal')
result.resid.plot(ax=ax4, title='Residual')

plt.show()
```

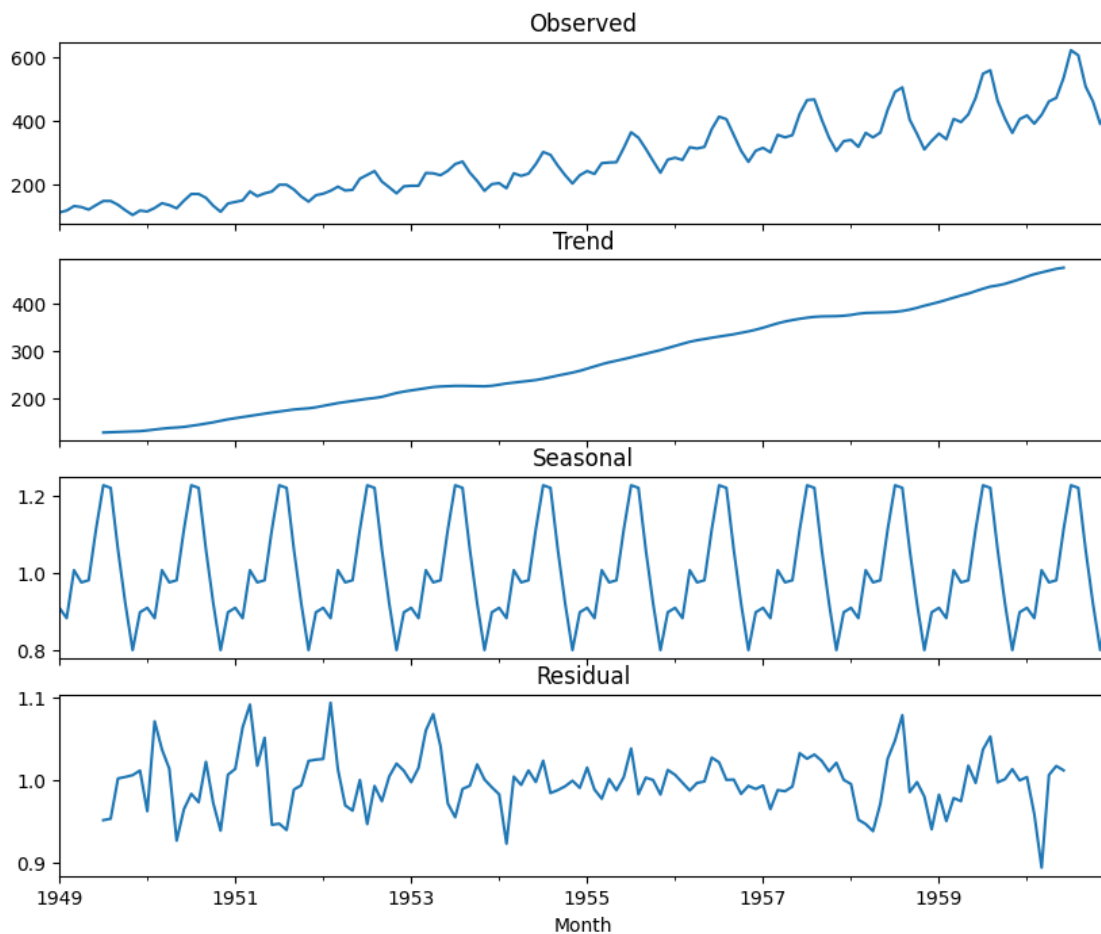


## 4.2 Multiplicative seasonal decomposition

```
[192]: result = seasonal_decompose(data['Passengers'], model='multiplicative')
fig, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(10, 8), sharex=True)

result.observed.plot(ax=ax1, title='Observed')
result.trend.plot(ax=ax2, title='Trend')
result.seasonal.plot(ax=ax3, title='Seasonal')
result.resid.plot(ax=ax4, title='Residual')

plt.show()
```



## 5 Build and evaluate time series forecast


### 5.1 Split time series data into training and test set

```
[193]: train_len = 120
train = data[:train_len] # first 120 months as the training set
test = data[train_len:] # last 24 months as the out-of-time test set
```

## 6 Auto Regressive methods

### 6.1 Stationarity vs non-stationary time series

```
[194]: data['Passengers'].plot(figsize=(12, 4))
plt.grid()
plt.legend(loc='best')
plt.title('Airline passenger traffic')
plt.show(block=False)
```



### 6.2 Augmented Dickey-Fuller (ADF) test

```
[195]: from statsmodels.tsa.stattools import adfuller
adf_test = adfuller(data['Passengers'])

print('ADF Statistic: %f' % adf_test[0])
print('Critical Values @ 0.05: %.2f' % adf_test[4]['5%'])
print('p-value: %f' % adf_test[1])
```

```
ADF Statistic: 0.815369
Critical Values @ 0.05: -2.88
p-value: 0.991880
```



### 6.3 Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test

```
[196]: from statsmodels.tsa.stattools import kpss
kpss_test = kpss(data['Passengers'])

print('KPSS Statistic: %f' % kpss_test[0])
print('Critical Values @ 0.05: %.2f' % kpss_test[3]['5%'])
print('p-value: %f' % kpss_test[1])
```

KPSS Statistic: 1.651312

Critical Values @ 0.05: 0.46

p-value: 0.010000


/tmp/ipykernel\_43/3602609379.py:2: InterpolationWarning: The test statistic is outside of the range of p-values available in the look-up table. The actual p-value is smaller than the p-value returned.

```
kpss_test = kpss(data['Passengers'])
```

### 6.4 Box Cox transformation to make variance constant


```
[197]: from scipy.stats import boxcox
data_boxcox = pd.Series(boxcox(data['Passengers'], lambda=0), index = data.index)

plt.figure(figsize=(12,4))
plt.grid()
plt.plot(data_boxcox, label='After Box Cox tranformation')
plt.legend(loc='best')
plt.title('After Box Cox transform')
plt.show()
```



## 6.5 Differencing to remove trend

```
[198]: data_boxcox_diff = pd.Series(data_boxcox - data_boxcox.shift(), data.index)
plt.figure(figsize=(12,4))
plt.grid()
plt.plot(data_boxcox_diff, label='After Box Cox tranformation and differencing')
plt.legend(loc='best')
plt.title('After Box Cox transform and differencing')
plt.show()
```



```
[199]: data_boxcox_diff.dropna(inplace=True)
data_boxcox_diff.tail()
```

```
[199]: Month
1960-08-01    -0.026060
1960-09-01    -0.176399
1960-10-01    -0.097083
1960-11-01    -0.167251
1960-12-01     0.102279
dtype: float64
```

## 6.6 Augmented Dickey-Fuller (ADF) test

```
[200]: adf_test = adfuller(data_boxcox_diff)

print('ADF Statistic: %f' % adf_test[0])
print('Critical Values @ 0.05: %.2f' % adf_test[4]['5%'])
print('p-value: %f' % adf_test[1])
```

```
ADF Statistic: -2.717131
Critical Values @ 0.05: -2.88
p-value: 0.071121
```

## 6.7 Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test

```
[201]: kpss_test = kpss(data_boxcox_diff)

print('KPSS Statistic: %f' % kpss_test[0])
print('Critical Values @ 0.05: %.2f' % kpss_test[3]['5%'])
print('p-value: %f' % kpss_test[1])
```

KPSS Statistic: 0.038304

Critical Values @ 0.05: 0.46


p-value: 0.100000

/tmp/ipykernel\_43/3639712988.py:1: InterpolationWarning: The test statistic is outside of the range of p-values available in the look-up table. The actual p-value is greater than the p-value returned.

```
kpss_test = kpss(data_boxcox_diff)
```


## 6.8 Autocorrelation function (ACF)

```
[202]: from statsmodels.graphics.tsaplots import plot_acf
plt.figure(figsize=(12,4))
plt.grid()
plot_acf(data_boxcox_diff, ax=plt.gca(), lags = 30)
plt.show()
```



## 6.9 Partial autocorrelation function (PACF)

```
[203]: from statsmodels.graphics.tsaplots import plot_pacf
plt.figure(figsize=(12,4))
plt.grid()
plot_pacf(data_boxcox_diff, ax=plt.gca(), lags = 30)
plt.show()
```



```
[204]: train_data_boxcox = data_boxcox[:train_len]
test_data_boxcox = data_boxcox[train_len:]
train_data_boxcox_diff = data_boxcox_diff[:train_len-1]
test_data_boxcox_diff = data_boxcox_diff[train_len-1:]
```

```
[205]: train_data_boxcox_diff[:10]
```

```
[205]: Month
1949-02-01    0.052186
1949-03-01    0.112117
1949-04-01   -0.022990
1949-05-01   -0.064022
1949-06-01    0.109484
1949-07-01    0.091937
1949-08-01    0.000000
1949-09-01   -0.084557
1949-10-01   -0.133531
1949-11-01   -0.134733
dtype: float64
```

## 7 Auto regression method (AR)

```
[206]: from statsmodels.tsa.arima.model import ARIMA
model = ARIMA(train_data_boxcox_diff, order=(1, 0, 0))
model_fit = model.fit()
print(model_fit.params)
```

```
const      0.009473
ar.L1      0.182911
sigma2     0.010733
dtype: float64
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```


```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

## 7.1 Recover original time series

```
[207]: y_hat_ar = data_boxcox_diff.copy()
y_hat_ar['ar_forecast_boxcox_diff'] = model_fit.predict(data_boxcox_diff.index.
    ↪min(), data_boxcox_diff.index.max())
y_hat_ar['ar_forecast_boxcox'] = y_hat_ar['ar_forecast_boxcox_diff'].cumsum()
y_hat_ar['ar_forecast_boxcox'] = y_hat_ar['ar_forecast_boxcox'].
    ↪add(data_boxcox[0])
y_hat_ar['ar_forecast'] = np.exp(y_hat_ar['ar_forecast_boxcox'])
```

```
[208]: plt.figure(figsize=(12,4))
plt.grid()
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_ar['ar_forecast'][test.index.min():], label='Auto regression_
    ↪forecast')
plt.legend(loc='best')
plt.title('Auto Regression Method')
plt.show()
```



```
[209]: rmse = np.sqrt(mean_squared_error(test['Passengers'],
    ↪ y_hat_ar['ar_forecast'][test.index.min():])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_ar['ar_forecast'][test.
    ↪ index.min():])/test['Passengers'])*100,2)
```

```
[210]: results = pd.DataFrame(columns=['Method', 'RMSE', 'MAPE'])
tempResults = pd.DataFrame({'Method':['Autoregressive (AR) method'], 'RMSE':
    ↪ [rmse], 'MAPE': [mape] })
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[210]:
```

	Method	RMSE	MAPE
0	Autoregressive (AR) method	93.42	13.72

## 8 Moving average method (MA)

```
[211]: model = ARIMA(train_data_boxcox_diff, order=(0, 0, 1))
model_fit = model.fit()
print(model_fit.params)
```

```
const      0.009523
ma.L1      0.258490
sigma2     0.010579
dtype: float64
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

### 8.1 Recover original time series

```
[212]: y_hat_ma = data_boxcox_diff.copy()
y_hat_ma['ma_forecast_boxcox_diff'] = model_fit.predict(data_boxcox_diff.index.
    ↪ min(), data_boxcox_diff.index.max())
y_hat_ma['ma_forecast_boxcox'] = y_hat_ma['ma_forecast_boxcox_diff'].cumsum()
```

```


y_hat_ma['ma_forecast_boxcox'] = y_hat_ma['ma_forecast_boxcox'].
    ↪add(data_boxcox[0])
y_hat_ma['ma_forecast'] = np.exp(y_hat_ma['ma_forecast_boxcox'])

```

```

[213]: plt.figure(figsize=(12,4))
plt.grid()
plt.plot(data['Passengers'][:train_len], label='Train')
plt.plot(data['Passengers'][train_len:], label='Test')
plt.plot(y_hat_ma['ma_forecast'][test.index.min():], label='Moving average_
    ↪forecast')
plt.legend(loc='best')
plt.title('Moving Average Method')
plt.show()

```



```

[214]: rmse = np.sqrt(mean_squared_error(test['Passengers'],
    ↪y_hat_ma['ma_forecast'][test.index.min():])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_ma['ma_forecast'][test.
    ↪index.min():])/test['Passengers'])*100,2)

```

```

[215]: tempResults = pd.DataFrame({'Method':['Moving Average (MA) method'], 'RMSE':
    ↪[rmse], 'MAPE': [mape] })
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results

```

```

[215]:

```

	Method	RMSE	MAPE
0	Autoregressive (AR) method	93.42	13.72
0	Moving Average (MA) method	91.61	13.40

## 9 Auto regression moving average method (ARMA)

```
[216]: model = ARIMA(train_data_boxcox_diff, order=(1, 0, 1))
model_fit = model.fit()
print(model_fit.params)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
```


```
const      0.009628
ar.L1      -0.581788
ma.L1       0.837584
sigma2      0.010129
dtype: float64
```

### 9.1 Recover original time series

```
[217]: y_hat_arma = data_boxcox_diff.copy()
y_hat_arma['arma_forecast_boxcox_diff'] = model_fit.predict(data_boxcox_diff.
    ↪index.min(), data_boxcox_diff.index.max())
y_hat_arma['arma_forecast_boxcox'] = y_hat_arma['arma_forecast_boxcox_diff'].
    ↪cumsum()
y_hat_arma['arma_forecast_boxcox'] = y_hat_arma['arma_forecast_boxcox'].
    ↪add(data_boxcox[0])
y_hat_arma['arma_forecast'] = np.exp(y_hat_arma['arma_forecast_boxcox'])
```

```
[218]: plt.figure(figsize=(12,4))
plt.grid()
plt.plot( data['Passengers'][:train_len-1], label='Train')
plt.plot(data['Passengers'][train_len-1:], label='Test')
plt.plot(y_hat_arma['arma_forecast'][test.index.min():], label='ARMA forecast')
plt.legend(loc='best')
plt.title('ARMA Method')
plt.show()
```





```
[219]: rmse = np.sqrt(mean_squared_error(test['Passengers'],
    ↪ y_hat_arma['arma_forecast'][train_len-1:])).round(2)
mape = np.round(np.mean(np.
    ↪ abs(test['Passengers']-y_hat_arma['arma_forecast'][train_len-1:])/
    ↪ test['Passengers'])*100,2)
```

```
[220]: tempResults = pd.DataFrame({'Method':['Autoregressive moving average (ARMA)
    ↪ method'], 'RMSE': [rmse], 'MAPE': [mape] })
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[220]:
```

	Method	RMSE	MAPE
0	Autoregressive (AR) method	93.42	13.72
0	Moving Average (MA) method	91.61	13.40
0	Autoregressive moving average (ARMA) method	88.74	12.81

## 10 Seasonal auto regressive integrated moving average (SARIMA)

```
[221]: from statsmodels.tsa.statespace.sarimax import SARIMAX

model = SARIMAX(train_data_boxcox, order=(1, 1, 1), seasonal_order=(1, 1, 1,
    ↪ 12))
model_fit = model.fit()
print(model_fit.params)
```

```
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
ValueWarning: No frequency information was provided, so inferred frequency MS
will be used.
```

```
self._init_dates(dates, freq)
/opt/conda/lib/python3.10/site-packages/statsmodels/tsa/base/tsa_model.py:473:
```

ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

\* \* \*

Machine precision = 2.220D-16

N = 5 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -1.61271D+00 |proj g|= 4.53288D+00

At iterate 5 f= -1.62802D+00 |proj g|= 1.10942D+00

At iterate 10 f= -1.63884D+00 |proj g|= 2.88919D-02

At iterate 15 f= -1.64511D+00 |proj g|= 1.67627D-01

At iterate 20 f= -1.64525D+00 |proj g|= 1.12474D-01

\* \* \*

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

\* \* \*

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
5	24	37	1	0	0	5.702D-03	-1.645D+00
F = -1.6452969897436860							

CONVERGENCE: REL\_REDUCTION\_OF\_F\_<=\_FACTR\*EPSMCH

ar.L1 -0.235263

ma.L1 -0.091737

ar.S.L12 -0.071385


ma.S.L12 -0.493694

sigma2 0.001403

dtype: float64

```
[222]: y_hat_sarima = data_boxcox_diff.copy()
y_hat_sarima['sarima_forecast_boxcox'] = model_fit.predict(data_boxcox_diff.
    ↪ index.min(), data_boxcox_diff.index.max())
y_hat_sarima['sarima_forecast'] = np.exp(y_hat_sarima['sarima_forecast_boxcox'])
```

```
[223]: plt.figure(figsize=(12,4))
plt.grid()
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_sarima['sarima_forecast'][test.index.min():], label='SARIMA_
    ↪ forecast')
plt.legend(loc='best')
plt.title('Seasonal autoregressive integrated moving average (SARIMA) method')
plt.show()
```



```
[224]: rmse = np.sqrt(mean_squared_error(test['Passengers'],
    ↪ y_hat_sarima['sarima_forecast'][test.index.min():])).round(2)
mape = np.round(np.mean(np.
    ↪ abs(test['Passengers']-y_hat_sarima['sarima_forecast'][test.index.min():])/
    ↪ test['Passengers'])*100,2)
```

```
[225]: tempResults = pd.DataFrame({'Method':['Seasonal autoregressive integrated_
    ↪ moving average (SARIMA) method'], 'RMSE': [rmse], 'MAPE': [mape] })
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[225]:
```

	Method	RMSE	MAPE
0	Autoregressive (AR) method	93.42	13.72
0	Moving Average (MA) method	91.61	13.40
0	Autoregressive moving average (ARMA) method	88.74	12.81
0	Seasonal autoregressive integrated moving aver...	44.71	8.85