twitter-sentiment-analysis

January 5, 2024

[49]:
```
#Importing Necessary Libraries
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from wordcloud import WordCloud, STOPWORDS
import re
import warnings
warnings.filterwarnings('ignore')
```

1 Data Loading

[50]:
```
train_df = pd.read_csv('/kaggle/input/twittersentimentdata/train.csv')
test_df = pd.read_csv('/kaggle/input/twittersentimentdata/test.csv')
```

[51]:
```
train_df.shape
```

[51]:
```
(31962, 3)
```

[52]:
```
train_df.duplicated().sum()
```

[52]:
```
0
```

[53]:
```
train_df.dtypes
```

[53]:
```
id        int64
label     int64
tweet    object
dtype: object
```

[54]:
```
train_df.isnull().sum()
```

[54]:
```
id       0
label    0
tweet    0
dtype: int64
```

[55]:
```
test_df.isnull().sum()
```

[55]:
```
id       0
tweet    0
dtype: int64
```

[56]:
```
# Plotting Word Clouds
stopwords = set(STOPWORDS)
stopwords.add('user')

def plot_wordcloud(tweets, title):
    wordcloud = WordCloud(width=800, height=800, background_color='white',␣
 ↪stopwords=stopwords, min_font_size=10).generate(tweets)
    plt.figure(figsize=(14, 6), facecolor=None)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.title(title, fontdict={'fontsize': 20})
    plt.show()
```

[57]:
```
negative_tweets = train_df['tweet'][train_df['label'] == 1].to_string()
positive_tweets = train_df['tweet'][train_df['label'] == 0].to_string()

plot_wordcloud(negative_tweets, 'Negative Tweets')
plot_wordcloud(positive_tweets, 'Positive Tweets')
```


Negative Tweets
racist
u
allahsoil
girl
stompi
temple vandalised
done
porn
still
believe
daily thank
call
please
good
guy
video
black amp
thlatest
will
â
might
woman
sea shepherd
libtard
sjw
police
think
racism
trump
say
obama
one
hate
sikh temple
amp
black
white
people
sex
hispanic amp
stop
comment
new
libtard
year
liberal
need
keep
man
amp
feel
women
sjw
muslim
thought factory
feminismiscancer
feminisміsterrorism
shepherd suppoers
polarisation tru
calgary wso
paladino
america
fuck
factory left
right polarisation
alt right
left right
protesting trump
carl paladino
racist antiraci

2 Feature Engineering

[58]:
```
# Feature Engineering
train_df_fe = train_df.copy()
train_df_fe['tweet_length'] = train_df_fe['tweet'].str.len()
train_df_fe['num_hashtags'] = train_df_fe['tweet'].str.count('#')
train_df_fe['num_exclamation_marks'] = train_df_fe['tweet'].str.count('!')
train_df_fe['num_question_marks'] = train_df_fe['tweet'].str.count('\?')
train_df_fe['total_tags'] = train_df_fe['tweet'].str.count('@')
train_df_fe['num_punctuations'] = train_df_fe['tweet'].str.count('[.,:;]')
train_df_fe['num_words'] = train_df_fe['tweet'].apply(lambda x: len(x.split()))
train_df_fe.head()
```

[58]:

	id	label	tweet	tweet_length	\
0	1	0	@user when a father is dysfunctional and is s…	102	
1	2	0	@user @user thanks for #lyft credit i can't us…	122	

```
2   3      0                                bihday your majesty            21
3   4      0  #model   i love u take with u all the time in …           86
4   5      0             factsguide: society now    #motivation            39

   num_hashtags  num_exclamation_marks  num_question_marks  total_tags  \
0             1                      0                   0           1
1             3                      0                   0           2
2             0                      0                   0           0
3             1                      3                   0           0
4             1                      0                   0           0

   num_punctuations  num_words
0                 1         18
1                 1         19
2                 0          3
3                 0         14
4                 1          4
```

[59]:
```
# Visualizing Relationship of Engineered Features with Sentiments
features = ['tweet_length', 'num_hashtags', 'num_exclamation_marks',␣
 ↪'num_question_marks', 'total_tags', 'num_punctuations', 'num_words']

# Check if train_df_fe has the expected columns
if set(features).issubset(train_df_fe.columns):
    plt.figure(figsize=(12, 16))
    colors = ['skyblue', 'salmon']

    for i, feature in enumerate(features, 1):
        plt.subplot(4, 2, i)
        sns.histplot(train_df_fe[train_df_fe.label == 0][feature],␣
 ↪label='Positive', kde=True, color=colors[0])
        sns.histplot(train_df_fe[train_df_fe.label == 1][feature],␣
 ↪label='Negative', kde=True, color=colors[1])
        plt.legend()

    plt.tight_layout()
    plt.show()
```

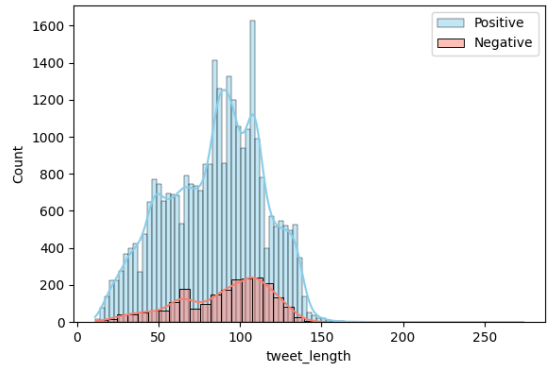
Positive
Negative
Count
tweet_length

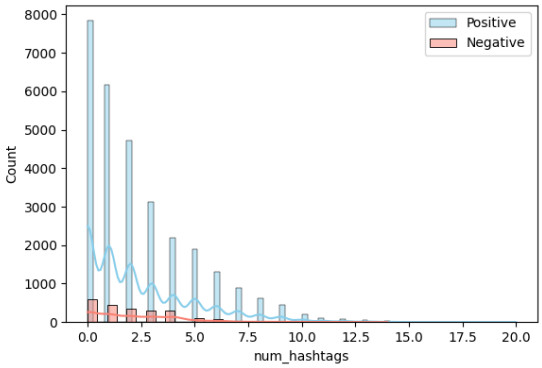
Positive
Negative
Count
num_hashtags

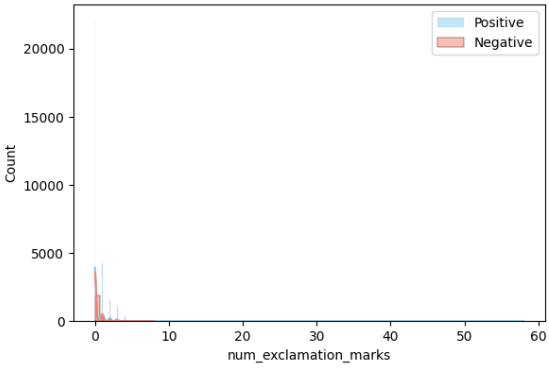
Positive
Negative
Count
num_exclamation_marks

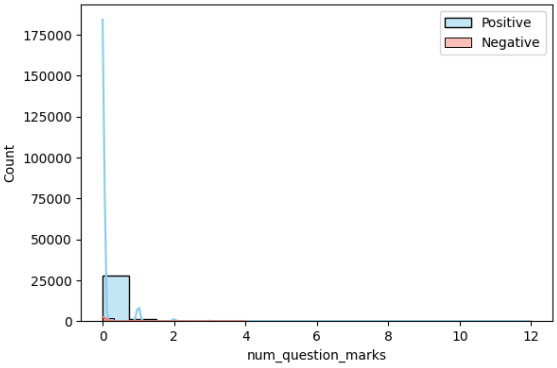
Positive
Negative
Count
num_question_marks

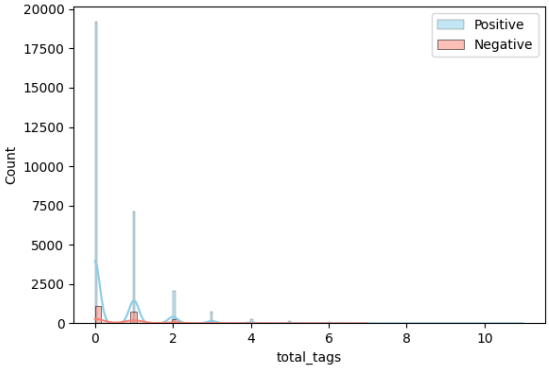
Positive
Negative
Count
total_tags

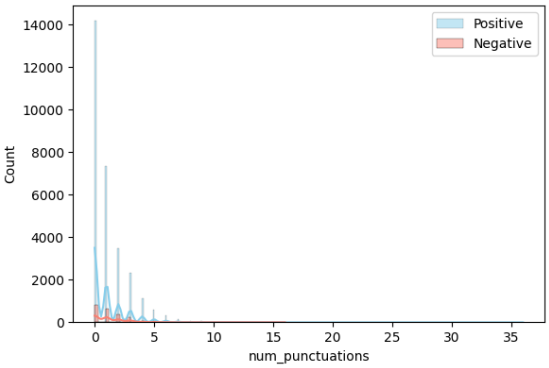
Positive
Negative
Count
num_punctuations

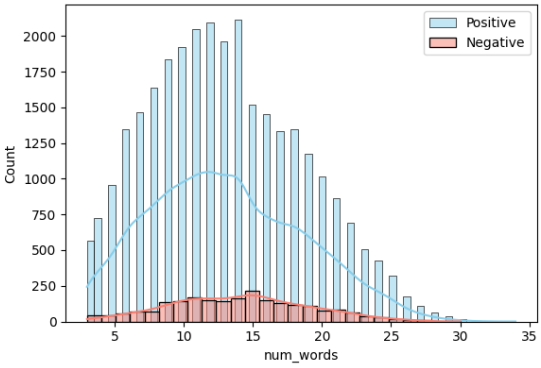
Positive
Negative
Count
num_words

3 Data Preprocessing

[60]:
```
# Data Preprocessing
X = train_df.drop(columns=['label'])
y = train_df['label']
test = test_df
```

[61]:
```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,␣
↪random_state=42)
```

3.1 Text Normalization

[62]:
```
def tokenize_and_clean(text):
    lowered = text.lower()
    cleaned = re.sub('@user', '', lowered)
    tokens = word_tokenize(cleaned)
    filtered_tokens = [token for token in tokens if re.match(r'\w{1,}', token)]
    stemmer = PorterStemmer()
    stems = [stemmer.stem(token) for token in filtered_tokens]
    return stems
```

3.2 Vectorization

[63]:
```
# TF-IDF Vectorization
tfidf_vectorizer = TfidfVectorizer(tokenizer=tokenize_and_clean,␣
↪stop_words='english')
X_train_tweets_tfidf = tfidf_vectorizer.fit_transform(X_train['tweet'])
X_test_tweets_tfidf = tfidf_vectorizer.transform(X_test['tweet'])
X_tweets_tfidf = tfidf_vectorizer.fit_transform(X['tweet'])
test_tweets_tfidf = tfidf_vectorizer.transform(test['tweet'])
```

3.3 SMOTE

[64]:
```
# Class Imbalance Check Before SMOTE
plt.figure(figsize=(12, 6))

# Colors for the pie charts
colors_before_smote = ['#66b3ff', '#99ff99']
colors_after_smote = ['#ff9999', '#66b3ff']

labels = ['Label 0 (Positive Tweets)', 'Label 1 (Negative Tweets)']

# Plotting before SMOTE
plt.subplot(1, 2, 1)
plt.pie(y_train.value_counts(), labels=labels, autopct='%0.1f%%',␣
↪colors=colors_before_smote)
```

```
plt.title('Class Distribution Before SMOTE')

# SMOTE to deal with the class imbalance
smote = SMOTE()
X_train_smote, y_train_smote = smote.fit_resample(X_train_tweets_tfidf, y_train.
 ↪values)

# Plotting after SMOTE
plt.subplot(1, 2, 2)
plt.pie(pd.value_counts(y_train_smote), labels=labels, autopct='%0.1f%%',␣
 ↪colors=colors_after_smote)
plt.title('Class Distribution After SMOTE')
plt.show()
```

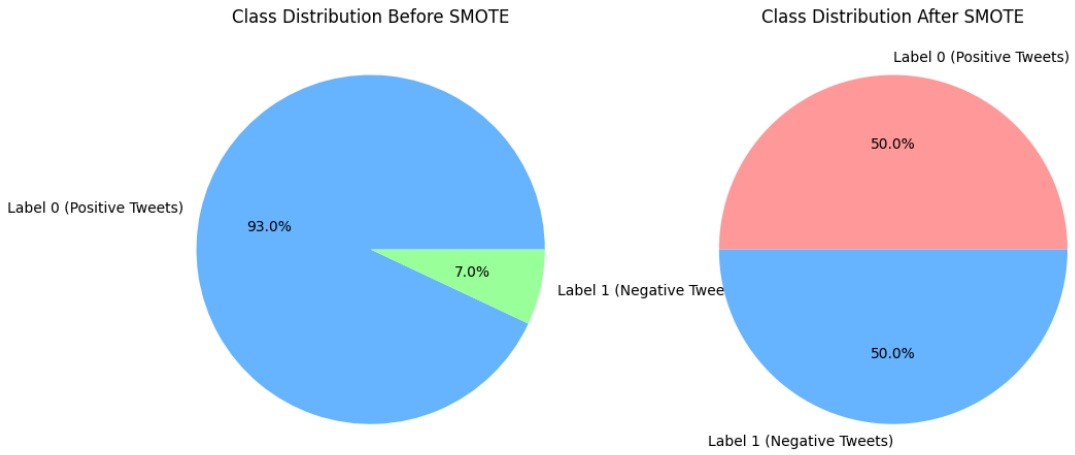

4 ML Model

[65]:
```
# Functions to print scores
def training_scores(y_act, y_pred):
    acc = round(accuracy_score(y_act, y_pred), 3)
    f1 = round(f1_score(y_act, y_pred), 3)
    print(f'Training Scores: Accuracy={acc}, F1-Score={f1}')

def validation_scores(y_act, y_pred):
    acc = round(accuracy_score(y_act, y_pred), 3)
    f1 = round(f1_score(y_act, y_pred), 3)
    print(f'Validation Scores: Accuracy={acc}, F1-Score={f1}')
```

[66]:
```
# Machine Learning Modeling
def train_and_evaluate(model, X_train, y_train, X_test, y_test):
    model.fit(X_train, y_train)
```

```
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)

    training_scores(y_train, y_train_pred)
    validation_scores(y_test, y_test_pred)
```

[67]:
```
# Logistic Regression
lr = LogisticRegression()
train_and_evaluate(lr, X_train_smote, y_train_smote, X_test_tweets_tfidf,␣
 ↪y_test)
```

```
Training Scores: Accuracy=0.974, F1-Score=0.975
Validation Scores: Accuracy=0.924, F1-Score=0.601
```

[68]:
```
# Naive Bayes Classifier
mnb = MultinomialNB()
train_and_evaluate(mnb, X_train_smote, y_train_smote, X_test_tweets_tfidf,␣
 ↪y_test)
```

```
Training Scores: Accuracy=0.966, F1-Score=0.967
Validation Scores: Accuracy=0.921, F1-Score=0.609
```

[69]:
```
# Random Forest Classifier
rf = RandomForestClassifier()
train_and_evaluate(rf, X_train_smote, y_train_smote, X_test_tweets_tfidf,␣
 ↪y_test)
```

```
Training Scores: Accuracy=1.0, F1-Score=1.0
Validation Scores: Accuracy=0.955, F1-Score=0.648
```

[70]:
```
# Extreme Gradient Boosting Classifier
xgb = XGBClassifier(objective='binary:logistic', eval_metric='logloss')
train_and_evaluate(xgb, X_train_smote, y_train_smote, X_test_tweets_tfidf,␣
 ↪y_test)
```

```
Training Scores: Accuracy=0.941, F1-Score=0.938
Validation Scores: Accuracy=0.942, F1-Score=0.597
```

5 Hyperparameter Tuning

[71]:
```
rf_tuned = RandomForestClassifier(criterion='entropy',
                                  max_samples=0.8,
                                  min_samples_split=10,
                                  random_state=0)
train_and_evaluate(rf_tuned, X_train_smote, y_train_smote, X_test_tweets_tfidf,␣
 ↪y_test)
```

```
Training Scores: Accuracy=0.999, F1-Score=0.999
Validation Scores: Accuracy=0.957, F1-Score=0.682
```

[72]:
```
xgb_tuned = XGBClassifier(objective='binary:logistic',
                          eval_metric='logloss',
                          learning_rate=0.8,
                          max_depth=20,
                          gamma=0.6,
                          reg_lambda=0.1,
                          reg_alpha=0.1)
train_and_evaluate(xgb_tuned, X_train_smote, y_train_smote,␣
 ↪X_test_tweets_tfidf, y_test)
```

```
Training Scores: Accuracy=0.998, F1-Score=0.998
Validation Scores: Accuracy=0.954, F1-Score=0.65
```

Your support with an **upvote** would be greatly valued.